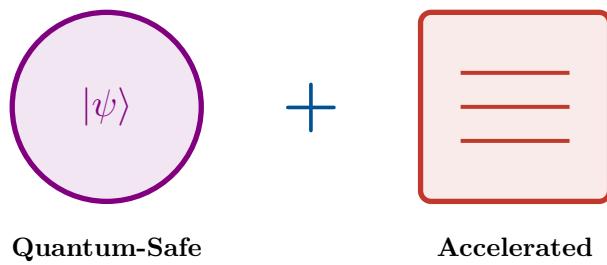




Post-Quantum Cryptography & Hardware Acceleration

GPU and FPGA-Based Compression
A Comprehensive Guide for Future-Proof Computing



Abdessamad JAOUAD

M2 Big Data & IoT
ENSAM Casablanca

January 3, 2026

Contents

| | | |
|----------|---|-----------|
| 1 | Introduction | 3 |
| 1.1 | The Quantum Threat | 3 |
| 1.1.1 | Timeline Estimates | 3 |
| 1.2 | The Performance Challenge | 3 |
| 1.3 | Hardware Acceleration Solutions | 3 |
| 1.4 | Document Overview | 4 |
| 1.5 | Why This Matters | 4 |
| 2 | Post-Quantum Cryptography Fundamentals | 5 |
| 2.1 | What is Post-Quantum Cryptography? | 5 |
| 2.2 | Why Current Cryptography Fails | 5 |
| 2.2.1 | Shor's Algorithm | 5 |
| 2.2.2 | Complexity Comparison | 5 |
| 2.3 | Grover's Algorithm | 6 |
| 2.4 | Mathematical Foundations of PQC | 6 |
| 2.5 | Security Levels | 6 |
| 3 | PQC Algorithm Families | 7 |
| 3.1 | Overview of Algorithm Families | 7 |
| 3.2 | Lattice-Based Cryptography | 7 |
| 3.2.1 | What is a Lattice? | 7 |
| 3.2.2 | Hard Lattice Problems | 7 |
| 3.2.3 | LWE Problem Explained | 7 |
| 3.3 | Hash-Based Signatures | 8 |
| 3.3.1 | One-Time Signatures (OTS) | 8 |
| 3.3.2 | Merkle Trees | 8 |
| 3.4 | Code-Based Cryptography | 8 |
| 3.4.1 | McEliece Cryptosystem (1978) | 8 |
| 3.5 | Algorithm Family Comparison | 9 |
| 4 | NIST PQC Standards | 10 |
| 4.1 | Standardization Timeline | 10 |
| 4.2 | The Four Standardized Algorithms | 10 |
| 4.3 | ML-KEM (Kyber) | 10 |
| 4.3.1 | How ML-KEM Works | 11 |
| 4.4 | ML-DSA (Dilithium) | 11 |
| 4.5 | SLH-DSA (SPHINCS+) | 11 |
| 4.6 | Comparison Summary | 11 |

| | |
|--|-----------|
| 5 GPU Architecture and CUDA | 13 |
| 5.1 GPU vs CPU Architecture | 13 |
| 5.2 CUDA Programming Model | 13 |
| 5.2.1 Key Concepts | 13 |
| 5.2.2 Thread Hierarchy | 14 |
| 5.3 CUDA Memory Model | 14 |
| 5.4 CUDA Kernel Example | 14 |
| 5.4.1 Execution Flow | 14 |
| 5.5 Performance Considerations | 15 |
| 5.6 CUDA Libraries for Compression | 15 |
| 6 OpenCL and Cross-Platform Computing | 16 |
| 6.1 Introduction to OpenCL | 16 |
| 6.1.1 Platform Model | 16 |
| 6.1.2 OpenCL Architecture Layers | 16 |
| 6.2 Memory Model | 16 |
| 6.3 OpenCL Kernel Programming | 17 |
| 6.3.1 Kernel Syntax | 17 |
| 6.3.2 Key OpenCL Functions | 17 |
| 6.4 CUDA vs OpenCL Comparison | 18 |
| 7 FPGA Architecture and Programming | 19 |
| 7.1 Introduction to FPGAs | 19 |
| 7.2 FPGA Architecture | 19 |
| 7.2.1 Core Components | 19 |
| 7.2.2 Look-Up Table (LUT) Structure | 19 |
| 7.3 FPGA vs GPU vs CPU | 20 |
| 7.4 FPGA Programming Methods | 20 |
| 7.4.1 Traditional HDL Flow | 20 |
| 7.4.2 High-Level Synthesis (HLS) | 20 |
| 7.5 Major FPGA Vendors | 21 |
| 8 Hardware-Accelerated Compression | 22 |
| 8.1 Why Hardware Acceleration? | 22 |
| 8.2 GPU Compression Techniques | 22 |
| 8.2.1 Parallel LZ77 | 22 |
| 8.2.2 Parallel Huffman Coding | 23 |
| 8.3 FPGA Compression Implementations | 23 |
| 8.3.1 Streaming LZ77 Architecture | 23 |
| 8.3.2 Commercial FPGA Compression IP | 23 |
| 8.4 Hybrid CPU-GPU-FPGA Systems | 23 |
| 8.5 PQC Hardware Acceleration | 24 |
| 8.5.1 PQC Operations on Hardware | 24 |
| 8.5.2 Number Theoretic Transform (NTT) | 24 |
| 9 Conclusion and Future Directions | 25 |
| 9.1 Key Takeaways | 25 |
| 9.2 Future Trends | 25 |
| 9.2.1 Emerging Technologies | 25 |
| 9.3 Recommendations | 26 |
| 9.3.1 Implementation Roadmap | 26 |
| 9.4 Summary Table | 26 |

Chapter 1

Introduction

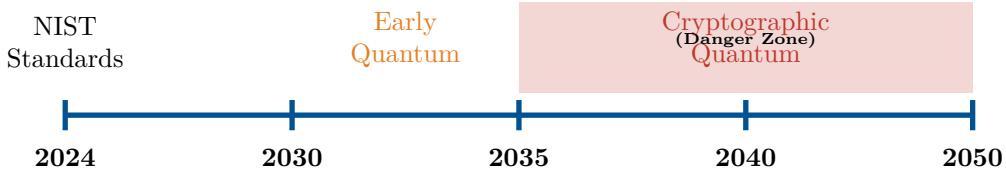
1.1 The Quantum Threat

The advent of **quantum computing** poses an existential threat to modern cryptography. Algorithms like **Shor's algorithm** can break RSA, ECC, and other public-key systems in polynomial time.

Harvest Now Decrypt Later Attack

Adversaries are already collecting encrypted data today, planning to decrypt it once quantum computers become available. This makes post-quantum migration **urgent**.

1.1.1 Timeline Estimates



1.2 The Performance Challenge

Post-quantum algorithms have significantly **larger key sizes** and **slower operations** than classical cryptography. Hardware acceleration becomes essential.

Table 1.1: Key Size Comparison: Classical vs. Post-Quantum

| Algorithm | Type | Public Key | Private Key |
|--------------|---------------|------------|-------------|
| RSA-2048 | Classical | 256 B | 1.2 KB |
| ECDSA P-256 | Classical | 64 B | 32 B |
| ML-KEM-768 | PQC (Lattice) | 1,184 B | 2,400 B |
| ML-DSA-65 | PQC (Lattice) | 1,952 B | 4,032 B |
| SLH-DSA-128s | PQC (Hash) | 32 B | 64 B |

1.3 Hardware Acceleration Solutions

Two main approaches exist for accelerating cryptographic and compression workloads:

CUDA / OpenCL**VHDL / Verilog****GPU**

Massively Parallel
Thousands of cores
High throughput

FPGA

Custom Logic
Low latency
Energy efficient

1.4 Document Overview

This guide covers two critical future trends in computing:

1. **Post-Quantum Cryptography (PQC)** – Cryptographic algorithms resistant to quantum attacks
2. **Hardware Acceleration** – GPU and FPGA-based acceleration for compression and cryptography

Chapter Roadmap

- **Chapter 2:** Post-Quantum Cryptography Fundamentals
- **Chapter 3:** PQC Algorithm Families
- **Chapter 4:** NIST PQC Standards
- **Chapter 5:** GPU Architecture & CUDA
- **Chapter 6:** OpenCL & Cross-Platform Computing
- **Chapter 7:** FPGA Architecture & Programming
- **Chapter 8:** Hardware-Accelerated Compression
- **Chapter 9:** Conclusion & Future Directions

1.5 Why This Matters

Critical Takeaway

The intersection of **post-quantum security** and **hardware acceleration** is not optional—it's essential for any system that needs to remain secure and performant in the coming decades.

Chapter 2

Post-Quantum Cryptography Fundamentals

2.1 What is Post-Quantum Cryptography?

Post-Quantum Cryptography (PQC) refers to cryptographic algorithms designed to resist attacks from both classical and quantum computers.

Key Distinction

PQC algorithms run on **classical computers**—they don't require quantum hardware. They're simply designed to resist quantum attacks.

2.2 Why Current Cryptography Fails

2.2.1 Shor's Algorithm

Shor's algorithm (1994) can efficiently solve two hard problems:

Quantum Vulnerable

Integer Factorization

$$N = p \times q$$

Breaks: RSA

Discrete Logarithm

$$g^x \equiv h \pmod{p}$$

Breaks: DH, ECDSA

2.2.2 Complexity Comparison

Table 2.1: Time Complexity: Classical vs Quantum

| Problem | Classical | Quantum (Shor) |
|----------------------------|------------------|----------------|
| Factoring n -bit integer | $O(e^{n^{1/3}})$ | $O(n^3)$ |
| Discrete Log | $O(e^{n^{1/3}})$ | $O(n^3)$ |

Exponential to Polynomial

Shor's algorithm reduces **exponential** problems to **polynomial** time—a catastrophic break for RSA and ECC.

2.3 Grover's Algorithm

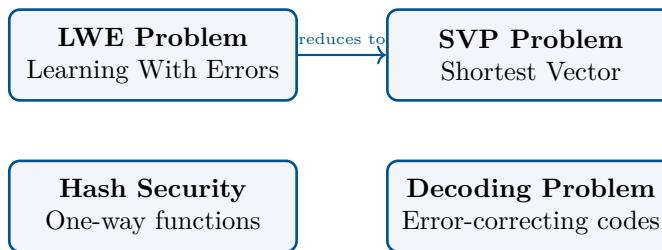
Grover's algorithm provides a **quadratic speedup** for searching unsorted databases:

- Classical search: $O(N)$
- Quantum search: $O(\sqrt{N})$

Impact on symmetric crypto: AES-128 becomes effectively AES-64. Solution: **double key sizes** (AES-256).

2.4 Mathematical Foundations of PQC

PQC relies on problems believed to be hard for *both* classical and quantum computers:



2.5 Security Levels

NIST defines 5 security levels based on equivalent classical security:

Table 2.2: NIST Security Levels

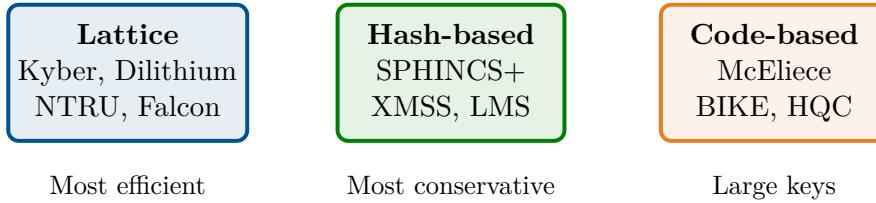
| Level | Equivalent Security | Bits |
|-------|---------------------|------|
| 1 | AES-128 key search | 128 |
| 2 | SHA-256 collision | 128 |
| 3 | AES-192 key search | 192 |
| 4 | SHA-384 collision | 192 |
| 5 | AES-256 key search | 256 |

Chapter 3

PQC Algorithm Families

Post-quantum algorithms are grouped into families based on their underlying mathematical problems.

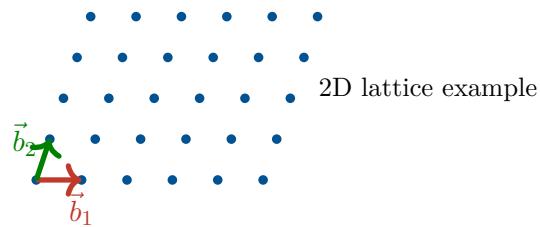
3.1 Overview of Algorithm Families



3.2 Lattice-Based Cryptography

3.2.1 What is a Lattice?

A **lattice** is a regular grid of points in n -dimensional space, generated by a set of basis vectors.



3.2.2 Hard Lattice Problems

1. **SVP** (Shortest Vector Problem): Find the shortest non-zero vector
2. **CVP** (Closest Vector Problem): Find the closest lattice point to a target
3. **LWE** (Learning With Errors): Distinguish noisy linear equations from random

3.2.3 LWE Problem Explained

Given samples of the form:

$$(\mathbf{a}_i, b_i = \langle \mathbf{a}_i, \mathbf{s} \rangle + e_i) \mod q$$

Where \mathbf{s} is secret and e_i is small noise. **Goal:** Recover \mathbf{s} .

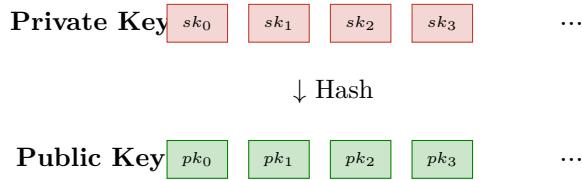
Why LWE is Hard

The noise e_i prevents simple linear algebra solutions. Without noise, solving is trivial. With noise, it becomes as hard as worst-case lattice problems.

3.3 Hash-Based Signatures

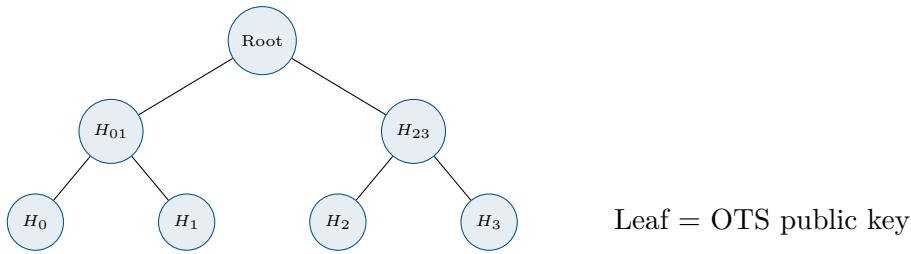
Hash-based signatures rely *only* on the security of hash functions—the most conservative assumption.

3.3.1 One-Time Signatures (OTS)



3.3.2 Merkle Trees

To sign multiple messages, OTS keys are organized in a **Merkle tree**:



3.4 Code-Based Cryptography

Based on the difficulty of decoding random linear codes.

3.4.1 McEliece Cryptosystem (1978)

1. Generate a decodable code with generator matrix G
2. Disguise it: $G' = SGP$ (scramble and permute)
3. Public key: G' Private key: S, G, P
4. Encrypt: $c = mG' + e$ (add errors)
5. Decrypt: Use private key to decode

Large Keys

McEliece public keys are ~ 1 MB at 128-bit security. This limits its use to specific applications.

3.5 Algorithm Family Comparison

Table 3.1: PQC Family Comparison

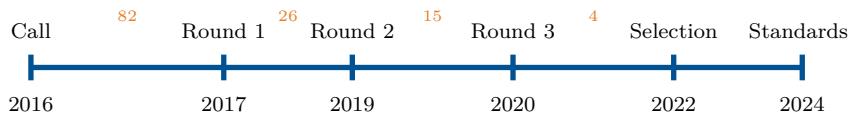
| Family | Key Size | Speed | Maturity | Best For |
|------------|-----------|-------|----------|--------------------|
| Lattice | Small | Fast | Medium | General use |
| Hash-based | Small/Med | Slow | High | Long-term security |
| Code-based | Large | Fast | High | Key encapsulation |

Chapter 4

NIST PQC Standards

In August 2024, NIST released the first post-quantum cryptography standards after an 8-year evaluation process.

4.1 Standardization Timeline



4.2 The Four Standardized Algorithms

FIPS 203: ML-KEM
(CRYSTALS-Kyber)
Key Encapsulation

FIPS 204: ML-DSA
(CRYSTALS-Dilithium)
Digital Signature

FIPS 205: SLH-DSA
(SPHINCS+)
Hash-based Signature

FIPS 206: FN-DSA
(Falcon)
Compact Signature

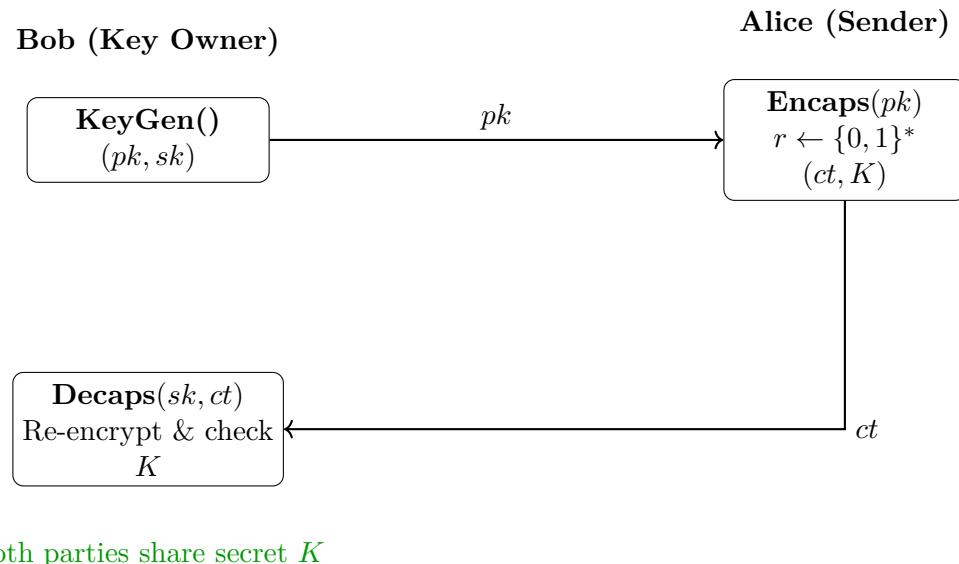
4.3 ML-KEM (Kyber)

Purpose: Key Encapsulation Mechanism (establishes shared secrets)

Table 4.1: ML-KEM Parameter Sets

| Variant | Security | Public Key | Ciphertext | Shared Secret |
|-------------|----------|------------|------------|---------------|
| ML-KEM-512 | Level 1 | 800 B | 768 B | 32 B |
| ML-KEM-768 | Level 3 | 1,184 B | 1,088 B | 32 B |
| ML-KEM-1024 | Level 5 | 1,568 B | 1,568 B | 32 B |

4.3.1 How ML-KEM Works



4.4 ML-DSA (Dilithium)

Purpose: Digital signatures (authentication)

Table 4.2: ML-DSA Parameter Sets

| Variant | Security | Public Key | Private Key | Signature |
|-----------|----------|------------|-------------|-----------|
| ML-DSA-44 | Level 2 | 1,312 B | 2,560 B | 2,420 B |
| ML-DSA-65 | Level 3 | 1,952 B | 4,032 B | 3,309 B |
| ML-DSA-87 | Level 5 | 2,592 B | 4,896 B | 4,627 B |

4.5 SLH-DSA (SPHINCS+)

Purpose: Conservative hash-based signatures

Why SPHINCS+?

SPHINCS+ relies **only** on hash function security. If lattice problems are broken, SPHINCS+ remains secure. It's the "backup plan."

Trade-off: Small keys (32-64 bytes) but large signatures (8-50 KB).

4.6 Comparison Summary

Table 4.3: NIST Standards Comparison (Level 3)

| Algorithm | Type | Pub Key | Priv Key | Sig/CT |
|--------------|-----------|---------|----------|----------|
| ML-KEM-768 | KEM | 1,184 B | 2,400 B | 1,088 B |
| ML-DSA-65 | Signature | 1,952 B | 4,032 B | 3,309 B |
| SLH-DSA-128f | Signature | 32 B | 64 B | 17,088 B |

Recommendation

For most applications: Use ML-KEM for key exchange and ML-DSA for signatures. Use SLH-DSA when conservative security assumptions are critical.

Chapter 5

GPU Architecture and CUDA

5.1 GPU vs CPU Architecture

GPUs and CPUs are designed for fundamentally different tasks:

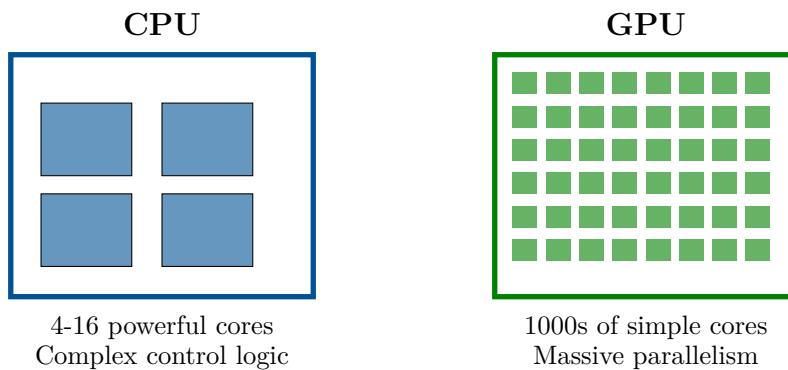


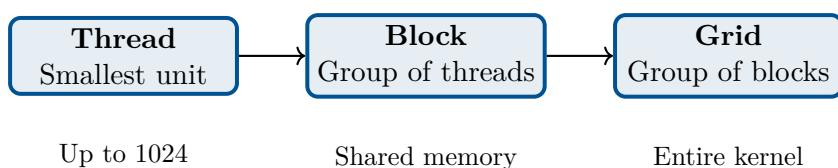
Table 5.1: CPU vs GPU Characteristics

| Feature | CPU | GPU |
|----------------|-------------------|----------------|
| Core count | 4-64 | 1,000-16,000 |
| Clock speed | 3-5 GHz | 1-2 GHz |
| Cache per core | Large (MB) | Small (KB) |
| Best for | Sequential tasks | Parallel tasks |
| Control flow | Complex branching | Simple SIMD |

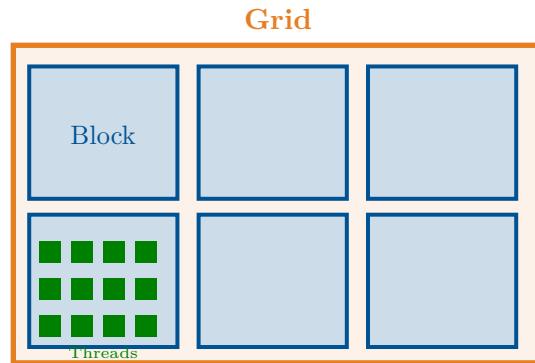
5.2 CUDA Programming Model

CUDA (Compute Unified Device Architecture) is NVIDIA's parallel computing platform.

5.2.1 Key Concepts



5.2.2 Thread Hierarchy



5.3 CUDA Memory Model

| | | |
|----|--|-------------|
| B | Registers (fastest, per-thread) | ~1 cycle |
| KB | Shared Memory (fast, per-block) | ~20 cycles |
| GB | Global Memory (slow, large) | ~400 cycles |

5.4 CUDA Kernel Example

A **kernel** is a function that runs on the GPU:

Vector Addition Kernel

```
--global__ void vectorAdd(float *a, float *b,
                         float *c, int n) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i < n) {
        c[i] = a[i] + b[i];
    }
}

// Launch: vectorAdd<<<numBlocks, blockSize>>>(...)
```

5.4.1 Execution Flow

1. Allocate GPU memory →
2. Copy data to GPU →
3. Execute kernel →
4. Copy results back

5.5 Performance Considerations

Optimization Tips

1. **Maximize parallelism:** Launch thousands of threads
2. **Minimize memory transfers:** CPU↔GPU is slow
3. **Use shared memory:** Cache frequently accessed data
4. **Coalesce memory access:** Adjacent threads access adjacent memory
5. **Avoid branch divergence:** Threads in a warp should follow same path

5.6 CUDA Libraries for Compression

NVIDIA provides optimized libraries:

Table 5.2: NVIDIA Compression Libraries

| Library | Algorithms | Throughput |
|---------|---------------------------------|----------------|
| nvCOMP | LZ4, Snappy, DEFLATE, zstd | Up to 500 GB/s |
| cuBLAS | Linear algebra (for transforms) | Optimized |
| cuFFT | FFT (for frequency domain) | Optimized |

nvCOMP Performance

nvCOMP can achieve **150×** speedup over CPU implementations for batch compression of small blocks.

Chapter 6

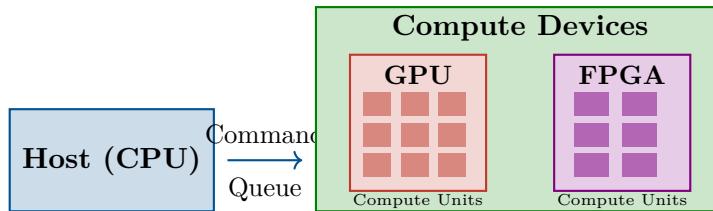
OpenCL and Cross-Platform Computing

6.1 Introduction to OpenCL

OpenCL (Open Computing Language) is an open standard for parallel programming across heterogeneous platforms—CPUs, GPUs, FPGAs, and DSPs from different vendors.

OpenCL provides **vendor-neutral** parallel computing, enabling the same code to run on AMD, Intel, NVIDIA, and ARM devices.

6.1.1 Platform Model



6.1.2 OpenCL Architecture Layers

1. **Platform Model:** Host + compute devices with compute units and processing elements
2. **Execution Model:** Kernels executed by work-items organized in work-groups
3. **Memory Model:** Global, local, constant, and private memory regions
4. **Programming Model:** Data-parallel and task-parallel execution

6.2 Memory Model

OpenCL defines four memory regions with different scope and performance characteristics:

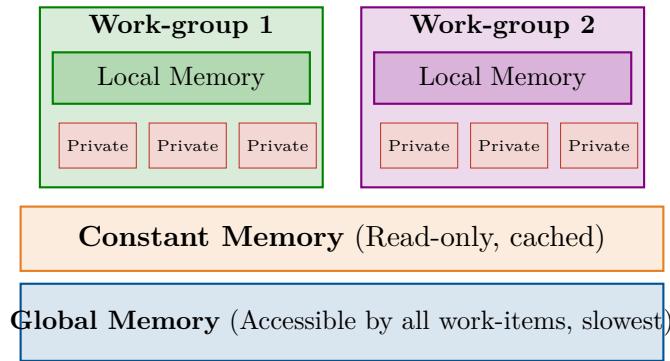


Table 6.1: OpenCL Memory Regions

| Memory | Scope | Speed | Size |
|----------|------------------|---------------|----------------|
| Global | All work-items | Slow | Large (GBs) |
| Constant | All (read-only) | Fast (cached) | Limited (64KB) |
| Local | Work-group | Fast | Small (48KB) |
| Private | Single work-item | Fastest | Very small |

6.3 OpenCL Kernel Programming

6.3.1 Kernel Syntax

OpenCL kernels are written in OpenCL C, a subset of C99 with extensions:

Listing 6.1: OpenCL Vector Addition Kernel

```
--kernel void vector_add(
    __global const float* A,
    __global const float* B,
    __global float* C,
    const int N)
{
    int gid = get_global_id(0);
    if (gid < N) {
        C[gid] = A[gid] + B[gid];
    }
}
```

6.3.2 Key OpenCL Functions

| Function | Purpose |
|----------------------|-----------------------------------|
| get_global_id(dim) | Global work-item ID |
| get_local_id(dim) | Local ID within work-group |
| get_group_id(dim) | Work-group ID |
| get_global_size(dim) | Total number of work-items |
| get_local_size(dim) | Work-group size |
| barrier() | Synchronization within work-group |

6.4 CUDA vs OpenCL Comparison

Table 6.2: CUDA vs OpenCL Feature Comparison

| Aspect | CUDA | OpenCL |
|----------------|------------------------|----------------------------------|
| Vendor | NVIDIA only | Cross-platform |
| Performance | Highly optimized | Vendor-dependent |
| Terminology | Threads, Blocks, Grids | Work-items, Work-groups, NDRange |
| Memory | Shared, Global | Local, Global |
| Ecosystem | Mature libraries | Broader hardware support |
| Learning Curve | Easier | More complex |

While OpenCL offers portability, **CUDA typically achieves 10-30% better performance** on NVIDIA hardware due to vendor-specific optimizations.

Chapter 7

FPGA Architecture and Programming

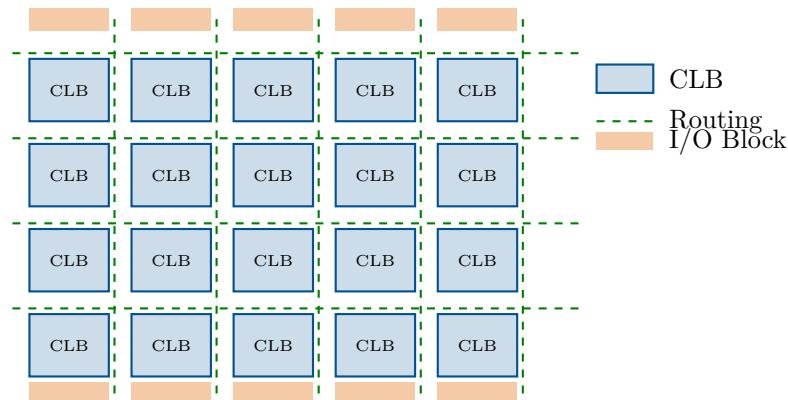
7.1 Introduction to FPGAs

Field-Programmable Gate Arrays (FPGAs) are integrated circuits that can be configured after manufacturing to implement custom digital logic.

Unlike GPUs that execute instructions, FPGAs implement algorithms **directly in hardware**, achieving extreme parallelism and energy efficiency.

7.2 FPGA Architecture

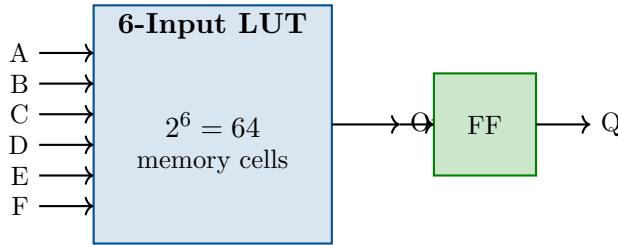
7.2.1 Core Components



- **Configurable Logic Blocks (CLBs):** Contain LUTs, flip-flops, and multiplexers
- **Programmable Interconnect:** Routing matrix connecting CLBs
- **I/O Blocks:** Interface with external devices
- **Block RAM (BRAM):** Dedicated memory blocks
- **DSP Slices:** Hardened multiply-accumulate units

7.2.2 Look-Up Table (LUT) Structure

A LUT implements any boolean function by storing the truth table in memory:



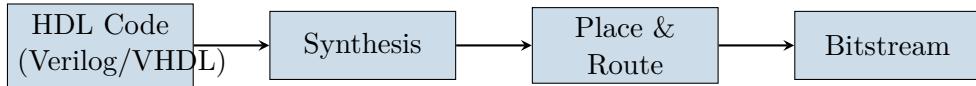
7.3 FPGA vs GPU vs CPU

Table 7.1: Hardware Platform Comparison

| Metric | CPU | GPU | FPGA |
|------------------|------------------|------------------|------------------|
| Parallelism | Low (8-64 cores) | High (thousands) | Extreme (custom) |
| Flexibility | Very high | High | Medium |
| Power Efficiency | Low | Medium | High |
| Latency | Medium | High | Very low |
| Development Time | Fast | Medium | Slow |
| Throughput | Low | Very high | High |

7.4 FPGA Programming Methods

7.4.1 Traditional HDL Flow



7.4.2 High-Level Synthesis (HLS)

Modern FPGAs support **High-Level Synthesis**, allowing C/C++/OpenCL code to be compiled to hardware:

Listing 7.1: HLS Example - Matrix Multiply

```

void matmul(float A[N][K], float B[K][M], float C[N][M]) {
    #pragma HLS INTERFACE m_axi port=A
    #pragma HLS INTERFACE m_axi port=B
    #pragma HLS INTERFACE m_axi port=C

    for (int i = 0; i < N; i++) {
        for (int j = 0; j < M; j++) {
            #pragma HLS PIPELINE II=1
            float sum = 0;
            for (int k = 0; k < K; k++) {
                sum += A[i][k] * B[k][j];
            }
            C[i][j] = sum;
        }
    }
}
  
```

HLS Pragmas

HLS pragmas guide the compiler to optimize for:

- PIPELINE: Enable pipelined execution
- UNROLL: Replicate loop iterations in hardware
- ARRAY_PARTITION: Distribute arrays across BRAMs

7.5 Major FPGA Vendors

| Vendor | Product Lines | Tools |
|--------------|----------------------------|------------------|
| AMD/Xilinx | Versal, UltraScale+, Artix | Vivado, Vitis |
| Intel/Altera | Agilex, Stratix, Cyclone | Quartus Prime |
| Lattice | CrossLink-NX, ECP5 | Radiant, Diamond |
| Microchip | PolarFire, SmartFusion | Libero SoC |

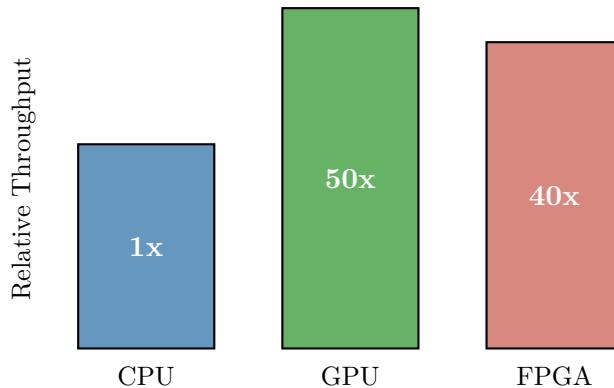
Chapter 8

Hardware-Accelerated Compression

8.1 Why Hardware Acceleration?

Data compression is computationally intensive. Hardware acceleration addresses:

1. **Throughput:** Process data faster than CPU alone
2. **Latency:** Reduce compression/decompression time
3. **Energy:** Lower power consumption per operation
4. **CPU Offloading:** Free CPU for other tasks

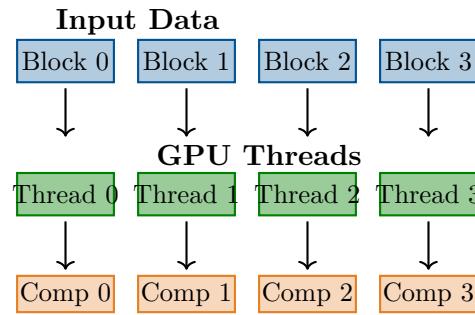


8.2 GPU Compression Techniques

8.2.1 Parallel LZ77

Traditional LZ77 is sequential (each match depends on previous output). GPU-optimized approaches:

- **Block-parallel:** Divide input into independent blocks
- **Hash-based matching:** Parallel hash table lookups
- **Batch processing:** Compress multiple streams simultaneously



8.2.2 Parallel Huffman Coding

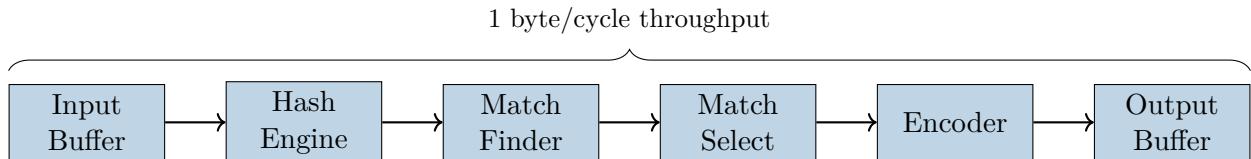
1. **Histogram:** Parallel frequency counting (atomic operations)
2. **Tree Building:** Sequential on CPU (small overhead)
3. **Encoding:** Parallel symbol-to-code mapping
4. **Bit Packing:** Parallel with careful synchronization

8.3 FPGA Compression Implementations

FPGAs excel at compression due to:

FPGAs implement compression as **streaming pipelines**—data flows through dedicated hardware stages with single-cycle latency per stage.

8.3.1 Streaming LZ77 Architecture



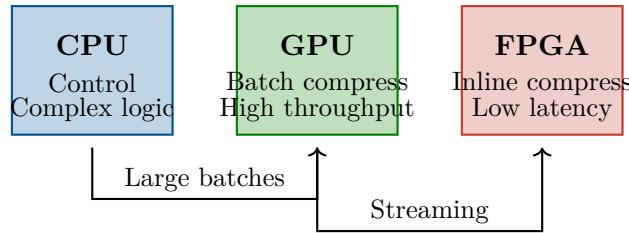
8.3.2 Commercial FPGA Compression IP

Table 8.1: FPGA Compression IP Cores

| Vendor | Algorithm | Throughput | Ratio |
|--------------|-----------|------------|--------|
| AMD/Xilinx | GZIP/ZSTD | 8 GB/s | 2.5-3x |
| Intel | GZIP | 10 GB/s | 2.5x |
| Atomic Rules | ZSTD | 100 Gbps | 3x |
| Eideticom | LZ4/ZSTD | 64 GB/s | 2-3x |

8.4 Hybrid CPU-GPU-FPGA Systems

Modern systems combine accelerators for optimal performance:



8.5 PQC Hardware Acceleration

Post-quantum cryptography is computationally expensive. Hardware acceleration is essential:

8.5.1 PQC Operations on Hardware

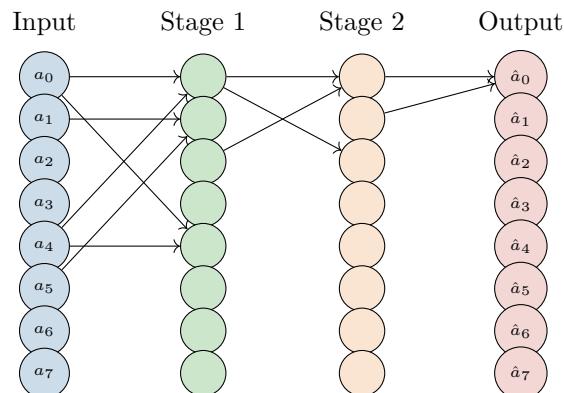
Table 8.2: PQC Hardware Acceleration Strategies

| Operation | GPU Approach | FPGA Approach |
|-------------------|------------------------|-----------------------|
| NTT (Lattice) | Parallel butterfly ops | Pipelined butterflies |
| Polynomial mult | Batch NTTs | Streaming NTT |
| Matrix operations | GEMM kernels | Systolic arrays |
| Hashing (SHAKE) | Parallel Keccak | Pipelined Keccak |
| Sampling | Parallel rejection | Hardware sampler |

Hardware-accelerated PQC can achieve **100-1000x** speedup over software implementations, making it practical for high-throughput applications.

8.5.2 Number Theoretic Transform (NTT)

NTT is the core operation in lattice-based cryptography. Similar to FFT but over finite fields:



NTT Parallelism

Each NTT stage has $n/2$ independent butterfly operations. GPUs process these in parallel; FPGAs pipeline them for streaming throughput.

Chapter 9

Conclusion and Future Directions

9.1 Key Takeaways

1. **Quantum Threat:** Shor's algorithm breaks RSA/ECC; Grover weakens symmetric crypto. Migration to PQC is urgent.
2. **NIST Standards:** ML-KEM, ML-DSA, SLH-DSA, and FN-DSA provide standardized post-quantum security.
3. **Performance Challenge:** PQC algorithms have larger keys and higher computational costs than classical cryptography.
4. **GPU Acceleration:** Massive parallelism enables batch processing; ideal for servers and data centers.
5. **FPGA Acceleration:** Custom pipelines offer low latency and energy efficiency; ideal for embedded and networking.
6. **Compression + PQC:** Hardware-accelerated compression reduces the bandwidth impact of larger PQC keys.

9.2 Future Trends



9.2.1 Emerging Technologies

- **PQC ASICs:** Custom chips for specific PQC algorithms
- **In-memory Computing:** Process data where it resides
- **Photonic Accelerators:** Optical computing for NTT operations
- **Quantum-Safe HSMs:** Hardware security modules with PQC support

9.3 Recommendations

Organizations should begin planning PQC migration **now**. The “Harvest Now, Decrypt Later” threat means sensitive data encrypted today may be compromised in the future.

9.3.1 Implementation Roadmap

1. **Inventory:** Catalog all cryptographic assets and dependencies
2. **Prioritize:** Focus on long-lived secrets and critical systems
3. **Test:** Evaluate PQC algorithms for performance and compatibility
4. **Hybrid:** Deploy hybrid classical/PQC solutions initially
5. **Monitor:** Track NIST guidance and quantum computing advances
6. **Accelerate:** Implement hardware acceleration where needed

9.4 Summary Table

Table 9.1: Technology Selection Guide

| Use Case | Recommended | Rationale |
|-----------------|----------------|-----------------------------------|
| Server/Cloud | GPU (CUDA) | High throughput, batch processing |
| Networking | FPGA | Low latency, inline processing |
| Embedded/IoT | FPGA | Power efficiency, small footprint |
| General Purpose | CPU + Software | Flexibility, ease of deployment |
| Hybrid Systems | GPU + FPGA | Combine throughput and latency |

The transition to post-quantum cryptography is not optional—it is a necessary evolution. Hardware acceleration will be essential to maintain performance in a post-quantum world.

“The quantum revolution in computing necessitates a quantum revolution in cryptography.”